

## [ Praticce Exercise ]

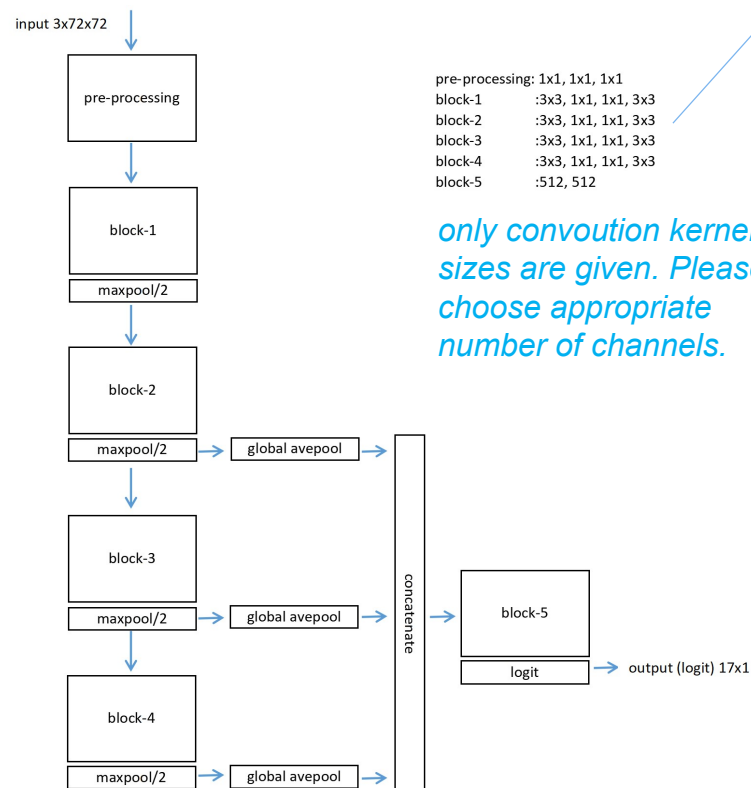
Given a pre-defined net, train and obtain validation and testing results.

### step.1

- resize the jpg train and test set from 256x256 to 72x72. Divide train set into 37479 training and 3000 validation samples.

### step.2

- implement the following net in pytorch.



```
class MyNet(nn.Module):  
  
    def __init__(self, in_shape=(3,72,72), num_classes=17):  
        super(MyNet, self).__init__()  
        in_channels, height, width = in_shape  
        stride=1  
  
        self.preprocess = nn.Sequential(  
            nn.Conv2d(in_channels, ??, kernel_size=1, stride=1, padding=0),  
            ??? #choose appropriate activation, etc  
            nn.Conv2d(??, ??, kernel_size=1, stride=1, padding=0),  
            ???  
            nn.Conv2d(??, ??, kernel_size=1, stride=1, padding=0),  
            ???  
        )  
  
        self.block1 = nn.Sequential(  
            ???  
            nn.MaxPool2d(kernel_size=2, stride=2),  
        )  
  
        self.block2 = nn.Sequential(  
            ???  
            nn.MaxPool2d(kernel_size=2, stride=2),  
        )  
        ...  
  
        self.block5 = nn.Sequential(  
            nn.Linear(???, 512),  
            ???  
            nn.Linear(512, 512),  
            ???  
        )  
  
        self.logit = nn.Linear(512, num_classes)  
  
    def forward(self, x):  
  
        out = self.preprocess(x)  
        out = self.block1(out)  
        out = self.block2(out)  
        ...  
        out = self.logit(out)  
        return out
```

- Hints:
  - choose appropriate number of channels
  - choose appropriate activation functions
  - you may need to add dropout, batch normalisation, etc
  - what is global average pooling?

<https://www.quora.com/What-is-global-average-pooling>

### step.3

- Implement training code to learn the net using stochastic gradient descent

```
optimizer = optim.SGD(net.parameters(), lr=?? , momentum=0.9, weight_decay=0.0005)

...

for epoch in range(num_epochs): # loop over the dataset multiple times

    ...
    for it, (images, labels, indices) in enumerate(train_loader, 0):

        logits = net(Variable(images))
        loss = criterion(logits, labels)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
```

- Hints:
  - choose appropriate loss function
  - choose appropriate batch size
  - choose appropriate number of epoches or stopping criteria
  - choose appropriate learning rates (may need learning rate scheduler)

- print your results as

epoch	iter	rate	train_loss	f2	valid_loss	f2	min
	???		0.157	0.855	0.172	0.820	0.5 min
			0.131	0.877	0.155	0.850	0.5 min
			0.159	0.855	0.135	0.870	0.5 min
			0.118	0.888	0.138	0.875	0.5 min
			0.120	0.891	0.125	0.890	0.5 min
			0.141	0.868	0.147	0.855	0.5 min
			0.109	0.892	0.125	0.888	0.5 min
			0.117	0.904	0.148	0.860	0.5 min
			0.142	0.853	0.132	0.878	0.5 min
			0.114	0.918	0.143	0.855	0.5 min
			0.108	0.913	0.137	0.868	0.5 min
			0.103	0.915	0.135	0.877	0.5 min
			0.106	0.907	0.107	0.90	
			0.126	0.889	0.112	0.90	
	???		0.105	0.895	0.111	0.90	
			0.110	0.907	0.110	0.90	
			0.102	0.913	0.120	0.89	
			0.102	0.912	0.120	0.88	
			0.145	0.876	0.147	0.86	
			0.097	0.927	0.122	0.88	
			0.127	0.885	0.106	0.91	
			0.107	0.908	0.096	0.92	
			0.097	0.925	0.095	0.91	
			0.103	0.904	0.093	0.921	0.5 min
			0.115	0.887	0.094	0.918	0.5 min
			0.109	0.913	0.093	0.920	0.5 min
			0.117	0.905	0.096	0.917	0.5 min
			0.089	0.924	0.094	0.920	0.5 min
			0.089	0.925	0.092	0.921	0.5 min
			0.081	0.942	0.093	0.922	0.5 min
			0.095	0.934	0.092	0.921	0.5 min
			0.087	0.936	0.091	0.923	0.5 min
			0.105	0.895	0.090	0.923	0.5 min
			0.088	0.913	0.090	0.922	0.5 min
			0.114	0.915	0.090	0.923	0.5 min
			0.104	0.895	0.090	0.922	0.5 min
	???		0.082	0.935	0.090	0.923	0.5 min

you need to implement f2 measure

<https://www.kaggle.com/c/planet-understanding-the-amazon-from-space#evaluation>

Submissions will be evaluated based on their mean  $F_2$  score. The F score, commonly used in information retrieval, measures accuracy using the precision  $p$  and recall  $r$ . Precision is the ratio of true positives  $tp$  to all predicted positives  $tp + fp$ . Recall is the ratio of true positives to all actual positives  $tp + fn$ . The  $F_2$  score is given by

$$(1 + \beta^2) \frac{pr}{\beta^2 p + r} \text{ where } p = \frac{tp}{tp + fp}, r = \frac{tp}{tp + fn}, \beta = 2.$$

Note that the  $F_2$  score weights recall higher than precision. The mean  $F_2$  score is formed by averaging the individual  $F_2$  scores for each row in the test set.

reference results on validation set is f2=0.923

#### **step.4**

- After training the net, apply it on the test set. Submit your results to the kaggle server and report your leader board f2 score.

#### **[ Passing mark ]**

f2 > 0.90 on validation set.

leader board score should be within 0.03 from your validation f2 score.

#### **[ Bounes ]**

f2 > 0.92 on validation set.

leader board score should be within 0.008 from your validation f2 score.